DeFi Micro■LLM — Tiered Plan of Attack

This document summarizes the three■tier strategy for building and integrating the DeFi micro■LLM, aligned with the NGF Stage■10/11 doctrine (Warp → Detect → Denoise). Each tier represents an increasing level of capability and integration.

Tier 0 — Baseline Deterministic Rails (✔ Secured)

• Stage■10 rails with matched filter + dual thresholds.

• Adapters: simple rule and stub mapper (ARC + DeFi).

• Residual traces and priors in place.

• Basic verifiers (ARC grid ops, DeFi invariants).

• End■to■end CLI runs are working.

Status: Complete — we already have this foundation secured.

Tier 1 — Micro■LLM on Synthetic Latents (Operational)

• Hybrid mapper + prior injection for stronger prompt→feature mapping.

• Stage■11 warp + detect + denoise rails wired against synthetic latent traces.

• Benchmarked successfully on ARC■like synthetic latents — demonstrates reasoning works.

• Micro■LLM exists: deterministic parser + denoiser operating on synthetic wells.

Status: Proven — benchmarks confirm deterministic reasoning on synthetic latents.

Tier 2 — Sidecar Integration with Real Latents (Aspirational)

• Real integration: extract pooled latents from an external LLM (e.g. GPT■2 sidecar).

• Replace synthetic latent generator with live embeddings for DeFi prompts.

• Micro■LLM rails (warp → detect → denoise) consume real latents to classify/sequence primitives.

• Bragging■point tier — showcases novel architecture, but high risk and currently unproven.

Status: Dream stage — ambitious future work, not required for MVP.

This tiered roadmap secures immediate deliverables (Tier 0 + Tier 1) while leaving Tier 2 as a forward■looking integration path. The MVP is viable at Tier 1, with Tier 2 reserved as a stretch goal.