Where We’re Going

The overarching direction is to build **micro reasoning AI machines** that apply across multiple task domains (ARC, math, coding, DeFi, etc.) by standardizing a **pipeline template**: 1. **Adapter** — map natural prompts + context into structured latent traces (domain■specific state schemas). 2. **Rails** — process those traces through NGF Stage■10/11 (Warp → Detect → Denoise). 3. **Verifier** — enforce domain■specific invariants (e.g., ARC grid transformations, DeFi LTV/HF checks). 4. **Pipelines** — glue components together into repeatable, testable reasoning flows. 5. **Sidecar Integration** — optional LLM hooks, allowing latent harvesting and hybrid pipelines. This roadmap ensures: - Deterministic reasoning (hallucination suppression by design). - Domain portability (adapters/verifiers swapped per category). - Compatibility with the existing **ngeodesic** package, allowing the micro■LLM to slot directly into the ecosystem. - Benchmarks and tests baked in from the start for reproducibility. In short, we are not just building a DeFi tool or ARC solver, but a **general micro reasoning template** that can be extended to any structured reasoning domain.