

DeFi Milestone Plan (Condensed to 5)

This document condenses the NGF 12-step process into 5 parsable milestones for DeFi, aligned with the Warp → Detect → Denoise doctrine described in the patents and article. Each milestone is script-driven (e.g., `defi_milestoneX.py`), validated, then folded into the `micro_llm` package with benchmarks and tests.

Milestone 0 — Baseline (■ done)

- Adapters: simple rule + stub mapper (ARC + DeFi).
- Residual traces + priors.
- Stage-10 rails (matched filter + dual thresholds).
- Basic verifiers (ARC grid ops, DeFi invariants).
- End-to-end CLI runs working.
- This is where we already are.

Milestone 1 — Hybrid Mapper + Prior Injection

- Integrate trained mapper (scikit-learn joblib stub).
- Confidence scoring + abstain wiring.
- Adapter residuals boosted with priors.
- Stage-10 rails respect priors when ordering.
- Produces a stronger front-end mapping from prompt → features → traces.

Milestone 2 — Stage-11 Warp + Detect

- Warp: funnel fit, PCA(3), curvature metrics.
- Detect: matched filtering with calibrated nulls (permutation / circular shifts).
- Output: stable well identification + ordered primitives.
- First NGF-native reasoning path, still without denoiser.

Milestone 3 — Stage-11 Denoise + Safety Guards

- Hybrid EMA+median smoothing.
- Confidence gates + noise floor rejection.
- Phantom-guard probes.
- Monte Carlo jitter averaging.
- Logging: SNR, phantom index, hallucination/omission rates.
- Yields deterministic suppression of phantom wells.

Milestone 4 — Consolidation + Benchmarks

- Fold into `micro_llm` package.
- ARC + DeFi benchmark suites (latent ARC, toy DeFi scenarios).
- Test harness for hallucination/abstain metrics.
- Push to GitHub, tag release.
- Demonstrates end-to-end deterministic micro-LLM reasoning.