

# Quickstart Flags & Modes — DeFi + ARC (Tier 1 & Tier 2)

This is the single, living reference for how to drive **micro-lm** via the **quickstart** entry point (CLI + Python). It unifies Tier-1 and Tier-2 semantics across **DeFi** and **ARC** and explains how flags interact.

---

## 0) TL;DR presets

### Tier-1 (mapper audit only, no WDD)

```
python3 -m micro_lm.cli.defi_quickstart "swap 2 ETH for USDC"
--rails stage11
--policy '{"mapper":{"model_path":".artifacts/
defi_mapper.joblib","confidence_threshold":0.70}}'
--verbose
```

### Tier-2 (WDD detector audit; mapper decides action)

```
python3 -m micro_lm.cli.defi_quickstart "deposit 10 ETH into aave"
--rails stage11
--policy '{"audit":{"backend":"wdd"},"mapper":
{"confidence_threshold":-1.0}}'
--verbose
```

### Tier-2 (WDD family; WDD supplies plan)

```
python3 -m micro_lm.cli.defi_quickstart "swap 2 ETH for USDC"
--rails stage11
--policy '{"audit":{"backend":"wdd","mode":"family","K":12,"template_width":
64,"z_abs":0.55,"keep_frac":0.70}}'
--verbose
```

### Edge block (example: DeFi LTV/HF/oracle)

```
python3 -m micro_lm.cli.defi_quickstart "withdraw 5 ETH"
--rails stage11
--policy '{"ltv_max":0.60,"mapper":{"confidence_threshold":-1.0}}'
```

```
--context '{"oracle":{"age_sec":5,"max_age_sec":30},"risk":{"hf":1.15}}'
--verbose
```

## 1) The single entry point

- CLI: `python3 -m micro_lm.cli.<domain>_quickstart "<prompt>" [--rails <stage>] [--policy JSON] [--context JSON] [--verbose]`
- Python: `from micro_lm.cli.<domain>_quickstart import quickstart → quickstart(prompt, policy=None, context=None, rails="stage11", T=180)`

**Uniform output contract** (both domains):

```
{
  "prompt": Str,
  "domain": "defi" | "arc",
  "rails": "stage11",
  "T": int,
  "top1": Str | null,
  "sequence": [Str],
  "plan": {"sequence": [Str]},
  "verify": {"ok": bool, "reason": Str, "tags": [Str]},
  "flags": { ... },
  "aux": { "stage11": { "wdd": { ... } }, ... },
  "wdd_summary": { "decision": "PASS" | "ABSTAIN" | null, "keep": [...], "sigma":
    int | null, "proto_w": int | null, "which_prior": Str | null, "note": Str | null },
  "det_hash": Str,
  "abstained": bool
}
```

Key invariants: - `plan.sequence` emits **canonical primitives** (e.g., `swap_asset`, `deposit_asset`). - If `verify.ok == false` (or reason contains a blocking keyword), **plan is cleared**: `plan.sequence = []`.

## 2) Policy schema (flags)

### 2.1 `policy.mapper` (Tier-1 audit; “action source” in all tiers)

Key	Type	Default	Meaning
<code>model_path</code>	str	null	Path to trained mapper (joblib).

Key	Type	Default	Meaning
<code>confidence_threshold</code>	float	0.70	Tier-1 <b>audit gate</b> ; if score < $\tau$ , mapper abstains. For Tier-2 runs, set to <code>-1.0</code> or use <code>skip_gate</code> to disable.
<code>skip_gate</code>	bool	false	(Optional toggle) If <code>true</code> , ignores <code>confidence_threshold</code> entirely. Useful for Tier-2 WDD.

**Notes** - The mapper is the **authoritative action chooser**; we always canonicalize its label to a primitive for `plan.sequence`. - For **Tier-2**, disable the gate (`confidence_threshold: -1.0` or `skip_gate: true`) so WDD is the audit.

## 2.2 `policy.audit` (Tier-2 audit)

Key	Type	Default	Meaning
<code>backend</code>	enum	null	Set to <code>"wdd"</code> to enable WDD.
<code>mode</code>	enum	<code>"detector"</code>	<code>"detector"</code> (audit only) or <code>"family"</code> (WDD returns <code>order</code> as plan).
<code>gate</code>	bool	false	If <code>true</code> , overwrite <code>verify.ok/reason</code> with WDD's decision (rare; detector is usually non-authoritative).
<code>K</code>	int	12	Family sampler size.
<code>template_width</code>	int	64	Family template width.
<code>z_abs</code>	float	0.55	Family z-score threshold.
<code>keep_frac</code>	float	0.70	Family keep fraction.
<code>overrides</code>	dict	<code>{}</code>	Family parameter overrides per prior.
<code>pca_prior</code>	str	null	Optional <code>.npz</code> for WDD PCA prior.
<code>debug</code>	bool	false	Extra debug in <code>aux.stage11.wdd.debug</code> ; enable <code>MICRO_LM_WDD_DEBUG=1</code> to print.

**Detector** behavior: - WDD attaches to `aux.stage11.wdd`, and `wdd_summary` mirrors essential fields. - `verify.reason` remains rails/local unless you set `audit.gate: true`.

**Family** behavior: - Short-circuits: returns `verify.reason = "wdd:family:<domain>"`, `flags.wdd_family = true`, and `plan.sequence = aux.wdd.order`.

## 2.3 Domain policy knobs

DeFi (examples): | Key | Type | Meaning | |---|---|---| | `ltv_max` | float | Max loan-to-value allowed. | | `hf_min` | float | Min health factor. | | `oracle.max_age_sec` | int | Price freshness requirement (in context). |

ARC (examples; align to your research checks): | Key | Type | Meaning | |---|---|---| | `must_cite` | bool | Block if no citation coverage. | | `rag.max_age_sec` | int | Retrieval freshness constraint. |

## 3) Environment variables

Var	Effect
<code>MICRO_LM_WDD_DEBUG=1</code>	Prints WDD detector debug lines (act, layer, sigma, prior, MF_peak).
<code>MICRO_LM_BENCHMARK_PLAN_OFF=1</code>	(Optional) Suppress <code>plan</code> for legacy benchmarks that only read <code>label</code> .

## 4) Semantics: who decides what?

1) **Mapper decides the action** (always). We canonicalize to `*_asset` (DeFi) or the ARC primitive set. 2) **WDD audits** the mapper's action. - Detector: advisory; metrics in `aux.stage11.wdd` + `wdd_summary`. - Family: authoritative planner when requested (policy `mode:"family"`). 3) **Verify = rails AND local domain verify**. - If local verify blocks (e.g., `ltv`, `hf`, `oracle` in DeFi), we **clear the plan** and set `verify.ok=false` with the domain reason.

`verify.reason` **cheat-sheet** - Rails shim only: `"shim:accept:stage-4"` (non-blocking default). - Detector annotated (optional): `"shim:accept:stage-4; wdd:pass"`. - Family: `"wdd:family:<domain>"`. - Local block (DeFi): contains `"ltv"`, `"hf"`, or `"oracle"`; plan is empty.

## 5) CLI patterns & gotchas

- **Quote JSON** with single quotes in zsh/macOS.
- `--use_wdd` (if present) is shorthand for `--policy '{"audit":{"backend":"wdd"}}'`.
- Disable Tier-1 gate in Tier-2 runs via `"mapper":{"confidence_threshold":-1.0}` or `"skip_gate":true`.
- Family mode parameters (`K`, `template_width`, `z_abs`, `keep_frac`) only apply in `mode:"family"`.

## 6) FAQ

**Q: Why do I still see `shim:accept:stage-4` when WDD is on?**

A: Detector mode is advisory by default. See `policy.audit.gate:true` to make WDD authoritative for `verify`.

**Q: How do I know WDD fired?**

A: `verify.tags` includes `audit:wdd` (detector), and `aux.stage11.wdd` + `wdd_summary` are populated. With `MICRO_LM_WDD_DEBUG=1` you'll see `[WDD]` lines.

**Q: Mapper vs `_infer_action`?**

A: WDD now prefers the **mapper-provided** canonical sequence; regex `_infer_action` is only a fallback if no sequence is given.

**Q: What clears the plan?**

A: Any `verify.ok=false` (e.g., `ltv/hf/oracle`), or explicit block reasons in the domain verifier.

---

## 7) Benchmark via quickstart (single surface)

Example DeFi Stage-8 run:

```
PYTHONWARNINGS="ignore::FutureWarning"
python3 scripts/tier2_benchmark.py
  --domains defi --runs 1
  --policy '{"mapper":{"model_path":".artifacts/
defi_mapper.joblib","confidence_threshold":0.35},"ltv_max":0.75,"hf_min":
1.0}'
  --outdir .artifacts
```

For Tier-2 WDD audits during the bench, disable the Tier-1 gate and add `audit.backend`:

```
PYTHONWARNINGS="ignore::FutureWarning"
python3 scripts/tier2_benchmark.py
  --domains defi --runs 1
  --policy '{"audit":{"backend":"wdd"},"mapper":
{"confidence_threshold":-1.0},"ltv_max":0.75,"hf_min":1.0}'
  --outdir .artifacts
```

## 8) Appendix — canonical primitive maps

### DeFi

swap	→ swap_asset
deposit	→ deposit_asset
withdraw	→ withdraw_asset
stake	→ stake_asset
unstake	→ unstake_asset
borrow	→ borrow_asset
repay	→ repay_asset
claim_rewards	→ claim_rewards_asset

### ARC (example; align to your templates)

classify	→ classify_text
extract	→ extract_spans
summarize	→ summarize_doc
qa rag	→ retrieve_and_answer

---

Keep this doc open while we integrate ARC. We'll add ARC's concrete policy keys (e.g., citation coverage) and examples as we wire them.